

Data Mining Exam

May 22, 2021

Kasper Rosenkrands

Aalborg University
Denmark



AALBORG UNIVERSITY
DENMARK

Clustering

Overview



- ▶ What is clustering?
- ▶ K-Means optimization problem and algorithm
- ▶ Implementation of the K-Means algorithm and an example
- ▶ Hierarchical Clustering (briefly)

Clustering

Introduction



Clustering is a way to categorize data to impose structure.

A use case is recommender systems (Amazon, Spotify, Netflix), where a user is recommended items that bought/listened to/watched by other users with similar interests.

Clustering

K-Means Optimization Problem

Given $D = (x_1, \dots, x_n)$ where $x_i \in \mathbb{R}^p$, $K \in \mathbb{N}$ and let C_1, \dots, C_K denote different groups of the x_i 's.

The K-Means algorithm tries to solve

$$\min_{C_1, \dots, C_K} \left\{ \sum_{k=1}^K W(C_k) \right\}, \quad (1)$$

where $W(C_k)$ denotes the **within cluster variation**, in other words the dissimilarity of the group.

The most common dissimilarity measure is the squared Euclidean distance

$$W(C_k) := \frac{1}{|C_k|} \sum_{i, i' \in C_k} \sum_{j=1}^p (x_{i,j} - x_{i',j})^2. \quad (2)$$

Clustering

K-Means Optimization Problem

If we by $\bar{x}_{k,j} = \frac{1}{|C_k|} \sum_{i \in C_k} x_{i,j}$ denote the mean value of the j 'th dimension in cluster k , it can be shown that

$$\frac{1}{|C_k|} \sum_{i, i' \in C_k} \sum_{j=1}^p (x_{i,j} - x_{i',j})^2 = 2 \sum_{i \in C_k} \sum_{j=1}^p (x_{i,j} - \bar{x}_{k,j})^2. \quad (3)$$

If we further note that $\bar{x}_{k,j} = \min_{\mu_k} \left\{ \sum_{i \in C_k} \sum_{j=1}^p (x_{i,j} - \mu_k)^2 \right\}$ this implies that the optimization problem in (1) can be rewritten as

$$\min_{C_1, \dots, C_k, \mu_1, \dots, \mu_k} \left\{ \sum_{k=1}^K \sum_{i \in C_k} \sum_{j=1}^p (x_{i,j} - \mu_k)^2 \right\}. \quad (4)$$

Clustering

K-Means Algorithm



The K-Means algorithm is now able to exploit the new formulation of the optimization problem and iteratively solve for $\{C_1, \dots, C_k\}$ and $\{\mu_1, \dots, \mu_k\}$.

This makes K-Means a greedy algorithm because, in each iteration it chooses optimal values for $\{C_1, \dots, C_k\}$ and $\{\mu_1, \dots, \mu_k\}$.

Convergence of the algorithm is therefore ensured, however we cannot guarantee it will find the global optimum.

Clustering

K-Means Algorithm



Algorithm 1: K-Means

```
1  Assign each observation to a cluster randomly
2  foreach Cluster do
3      |   Compute the centroid
4  foreach Observation do
5      |   Compute distance to all centroids
6      |   Assign to the closest
7  while Centroids have not changed since last iteration do
8      |   foreach Observation do
9          |   Compute distance to all centroids
10         |   Assign to the closest
11         foreach Cluster do
12             |   Compute the centroid
13 return Clusters
```

Clustering

An example of the K-Means algorithm




Figure: Iteration 01

Clustering

An example of the K-Means algorithm




Figure: Iteration 02

Clustering

An example of the K-Means algorithm




Figure: Iteration 03

Clustering

An example of the K-Means algorithm




Figure: Iteration 04

Clustering

An example of the K-Means algorithm




Figure: Iteration 05

Clustering

An example of the K-Means algorithm




Figure: Iteration 06

Clustering

An example of the K-Means algorithm




Figure: Iteration 07

Clustering

An example of the K-Means algorithm




Figure: Iteration 08

Clustering

An example of the K-Means algorithm




Figure: Iteration 09

Clustering

An example of the K-Means algorithm




Figure: Iteration 10

Clustering

An example of the K-Means algorithm




Figure: Iteration 11

Clustering

An example of the K-Means algorithm




Figure: Iteration 12

Clustering

An example of the K-Means algorithm




Figure: Iteration 13

Clustering

Hierarchical Clustering



Todo

- ▶ Introduction
- ▶ Type - Agglomerative vs Divise
- ▶ Pseudocode for algorithm, or just in words
- ▶ Visualization with dendogram
- ▶ Linkage types - (complete, single, average, centroid)

Shrinkage

Overview



Write equations and theory about ridge, lasso and elastic net

Find an example to compare the 3 methods on

Shrinkage

Variable Selection and Regularization



Because it is often cheaper to obtain multiple observations from a few samples than to obtain more samples, thus increasing the number of explanatory variables in a regression model, a linear regression model would be prone to increasing variance.

When extending linear regression to multiple explanatory variables the main objectives is:

- ▶ **Model Interpretability:** Models with fewer variables are often more easy to interpret results from and are therefore better to use for decision making.
- ▶ **Prediction Accuracy:** If by introducing some bias we are able to dramatically improve prediction accuracy, this would be worth considering (bias-variance tradeoff).

Shrinkage

Variable Selection and Regularization



There are multiple tools we can use in this pursuit

- ▶ **Subset Selection:** Works by fitting lots of models with different combinations of predictors. Then we can find out which variables are most related to the response and we can select these.
- ▶ **Dimensionality Reduction:** Works by projecting explanatory variables into a smaller dimensional space and use these projections as predictors.
- ▶ **Shrinkage Methods:** Works by fitting a model using all predictors while shrinking coefficients towards zero, to reduce variance. Some shrinkage methods can also perform variable selection by shrinking coefficients to exactly zero.

Shrinkage

Variable Selection and Regularization



This presentation will focus on shrinkage methods. The two main shrinkage methods are

- ▶ **Ridge Regression**
- ▶ **Lasso**

They both penalize the “size” of the estimated parameters, however they differ in the way they quantify the “size”. Ridge regression penalizes the ℓ_2 norm while Lasso penalizes the ℓ_1 norm.

Therefore we can also refer to the two methods as ℓ_1 - and ℓ_2 -regularizations.

Shrinkage

Ridge Regression



Write the equation that ridge minimizes

Shrinkage

Lasso



Write the equation that Lasso minimizes

Shrinkage

Elastic Net



The method is a compromise between Ridge Regression and Lasso as it minimizes

Write the equation that elastic net minimizes

Classification

Overview



Classification

Introduction

If the response variable is categorical (qualitative), i.e. it is of the form $y \in \{1, \dots, L\}$. Then a linear model of the form

$$y = \beta_0 + \beta_1 x_1 + \dots + \beta_p x_p,$$

is generally not a good approach to take as it could predict invalid values and for certain types of categorical data there might not be a clear ordering.

Therefore when dealing with categorical variable the aim of the model is to predict the probability that an observation belongs to a certain category, rather than the category itself.

Classification

Linear Discriminant Analysis (LDA)

Suppose we have a dataset with a response variable $y \in \{0, \dots, L\}$, explanatory variables x_1, \dots, x_p and that we would like to model

$$P(y = k|X) = \underbrace{\frac{P(y = k)P(X|y = k)}{P(X)}}_{\text{Bayes Theorem}} = \frac{\pi_k f_k(x)}{\sum_{i=1}^K f_i(x)}, \quad (5)$$

where $\pi_k = P(y = k)$ and $f_k(x) = P(X|y = k)$.

If we use the proportion of observations in the dataset that belong to class k as an estimate for π_k , then we just need to model $f_k(x) = P(X|y = k)$.

In other words we need to make some assumption on the distribution of $f_k(x)$.

Classification

Linear Discriminant Analysis (LDA)

The assumption made in LDA is that each $f_k(x)$ come from a multivariate normal distribution, i.e.

$$f_k(x) = \frac{1}{(2\pi)^{p/2} |\Sigma_k|^{1/2}} \exp \left\{ -\frac{1}{2} (x - \mu_k)^\top \Sigma_k^{-1} (x - \mu_k) \right\}. \quad (6)$$

Another assumption made in LDA is that $\Sigma_k = \Sigma \forall k \in \{1, \dots, K\}$, in other words all classes will have the same variance-covariance matrix.

Classification

Quadratic Discriminant Analysis (QDA)



Classification

Naive Bayes



Trees

Classification and Regression Trees (CART)



Trees

Bagging



Trees

Random Forest



Trees

Boosting



Support Vector Machines



Neural Networks

Perceptron



Backpropagation: For a given loss function L we look for $\frac{\partial L}{\partial w_i}$. We start with initial values for the weights, which we shall denote w_{old} . Then we update the weights by $w_{new} = w_{old} - \eta \frac{\partial L}{\partial w}$. One iteration is called an **epoch**.

